

The Role of Statistical Foundations in Machine Learning: Bias, Uncertainty, and the Limits of Predictive Models

Ilgiz Satkynov

Department of Mathematics

Physical Sciences, Mathematics and Engineering Division

De Anza College

June 16, 2026

Abstract

Artificial intelligence systems are frequently described as autonomous, data-driven, and objective, yet every predictive model rests on a foundation of classical statistical theory. This paper examines how core statistical concepts, including probability distributions, bias-variance tradeoff, and sampling theory, govern the behavior, accuracy, and limitations of modern machine learning systems. Through an analysis of regression-based and neural network-based models, including a small original simulation of the bias-variance decomposition, this paper argues that statistical bias in training data is not a peripheral engineering flaw but a structural consequence of how predictive models generalize from samples to populations. Case studies in hiring algorithms, facial recognition, and criminal risk assessment illustrate how mathematical properties of estimators translate into real social consequences. The paper concludes that statistical literacy, rather than purely computational power, is the critical skill needed to build, audit, and trust AI systems responsibly.

Table of Contents

1. Introduction	3
2. Statistical Foundations of Predictive Models	4
2.1 From Probability to Prediction	4
2.2 The Bias-Variance Tradeoff	6
2.3 Sampling, Representativeness, and the Population Problem	9
3. Case Studies: When Statistical Bias Becomes Social Bias	11
3.1 Hiring Algorithms	11
3.2 Facial Recognition and Subgroup Error Rates	11
3.3 Criminal Risk Assessment Tools	12
4. Statistical Remedies and Their Limits	13
4.1 Reweighting and Resampling	13
4.2 Fairness Constraints and Their Statistical Tradeoffs	14
4.3 The Limits of Post Hoc Correction	14
4.4 Statistical Auditing in Practice	15
5. Discussion: Why These Are Statistical Problems, Not Just Engineering Problems	15
6. Conclusion	16
References	18

1. Introduction

Machine learning is often presented to the public as a triumph of computer science: faster processors, larger datasets, and clever neural network architectures that can outperform humans at games, image recognition, and language generation. This framing, while not wrong, is incomplete. Underneath every machine learning model is a statistical engine, one built from probability theory, estimation, and inference that long predates the field of artificial intelligence itself. Logistic regression, the workhorse of binary classification, was developed by statisticians decades before “machine learning” was a common term. Neural networks, often treated as black boxes, are fundamentally nonlinear function approximators trained using statistical optimization methods to minimize expected error over a sample drawn from some unknown population.

This paper investigates the relationship between statistics and artificial intelligence from three angles. First, it reviews the statistical principles that underlie predictive modeling, focusing on the bias-variance tradeoff and the assumptions models make about the data they are trained on. Second, it analyzes how violations of these statistical assumptions, particularly unrepresentative sampling, produce biased outcomes in deployed AI systems. Third, it discusses why these issues are not simply “bugs” to be patched but are mathematically predictable consequences of how estimators behave on finite samples. The paper draws on published research in statistics, computer science, and the growing field of algorithmic fairness in order to support each of these claims.

The central argument of this paper is that statistical foundations should be treated as the primary lens through which AI systems are evaluated, not merely a technical detail subordinate to model architecture. A poorly specified statistical model trained on a perfectly clean dataset will still fail in predictable, quantifiable ways; understanding those ways requires statistical reasoning, not just programming skill.

The motivation for this paper comes from a broader gap often observed between how artificial intelligence is taught and how it is discussed publicly. Introductory machine learning courses frequently emphasize implementation, how to load a dataset, choose a library, and train a model, while treating the underlying statistical assumptions as background material to be covered briefly before moving on to applications. Public discussion of AI bias, meanwhile, often treats biased outcomes as a matter of corporate negligence or insufficient ethical oversight, without engaging with the precise statistical mechanisms that make biased outcomes a near-mathematical certainty under certain data conditions. This paper attempts to bridge that gap by walking through the statistical mechanisms in enough technical detail to be useful,

including a small original simulation, while remaining accessible to an undergraduate audience with a background in introductory statistics and computer science rather than a specialized machine learning theory course.

The remainder of the paper proceeds as follows. Section 2 develops the statistical foundations needed to discuss model behavior precisely, including probability calibration, the bias-variance decomposition, and the role of sampling theory in determining whether a model generalizes well. Section 3 applies these foundations to three documented cases of AI systems producing disparate outcomes across demographic groups. Section 4 reviews common statistical remedies, including reweighting and fairness constraints, along with their documented limitations. Section 5 discusses the broader implications of treating these issues as statistical rather than purely engineering problems, and Section 6 concludes.

2. Statistical Foundations of Predictive Models

2.1 From Probability to Prediction

At its core, a predictive model is an estimate of a conditional probability distribution: given some input data, what is the most likely outcome? In supervised learning, a model is trained on a sample of paired observations, inputs and known outcomes, and is asked to learn a function that generalizes to new, unseen inputs drawn from the same underlying population. This generalization step is fundamentally a statistical inference problem, no different in spirit from a pollster trying to predict an election outcome from a sample of voters.

The model's stated confidence in a prediction, such as the output of a logistic regression or the softmax layer of a neural network, is meant to represent a probability. Whether that probability is well-calibrated, meaning whether predicted probabilities actually match observed frequencies, is an empirical statistical question. A model that says an event will happen “90% of the time” should be correct approximately nine times out of ten when that prediction is made repeatedly. Many modern machine learning systems, particularly deep neural networks, are known to be poorly calibrated, often expressing higher confidence than their accuracy warrants (Guo et al., 2017), a finding that has been studied extensively in the statistics and machine learning literature on calibration (Dawid, 1982).

It is useful to recall the statistical machinery that produces these probability estimates in the first place. In logistic regression, parameters are typically chosen by maximum likelihood estimation: the model selects the coefficients that make the observed training labels as probable as possible under the assumed model form. In a neural network classifier, the analogous step is

minimizing cross-entropy loss, which is mathematically equivalent to maximum likelihood estimation under a categorical distribution assumption. Despite the very different computational machinery, gradient descent over millions of parameters versus closed-form or iteratively reweighted least squares for a handful of coefficients, both procedures are answering the same statistical question: which parameter values make the observed data least surprising? This shared foundation is precisely why statistical pathologies such as miscalibration, overconfidence, and sensitivity to sample composition appear in both simple and highly complex models.

A further statistical subtlety is the distinction between a model's discriminative performance, its ability to rank-order outcomes correctly, and its calibration, its ability to assign accurate probabilities to those outcomes. A model can have excellent discriminative ability, correctly identifying which applicants are more likely to default on a loan than others, while still being poorly calibrated, systematically overstating or understating the actual probability of default for every applicant. Metrics commonly reported in machine learning papers, such as the area under the receiver operating characteristic curve, measure discrimination but say little about calibration. A research-literate reader of an AI system's reported accuracy should therefore ask which of these two statistical properties, discrimination or calibration, is actually being claimed.




Figure 1. A reliability diagram comparing a well-calibrated model against an overconfident model. The well-calibrated model's points lie close to the diagonal, where predicted probability matches observed

frequency; the overconfident model systematically pushes predictions toward the extremes. Curves are illustrative, generated from simulated data rather than a specific published model.

It is worth pausing on why overconfidence specifically, rather than underconfidence, is the typical failure mode for large neural networks. One commonly cited explanation in the calibration literature is that the cross-entropy loss function used to train these models, combined with very high model capacity, rewards pushing predicted probabilities toward 0 or 1 even after the model has already learned to rank examples correctly, since doing so continues to reduce the loss on the training set without necessarily improving generalization. This is, again, a statistical phenomenon traceable to the interaction between an optimization objective and model capacity, not an arbitrary quirk of any particular software implementation.

2.2 The Bias-Variance Tradeoff

The bias-variance tradeoff is one of the most important statistical concepts for understanding why machine learning models succeed or fail. Formally, the expected prediction error of a model at a given point can be decomposed into three components: the square of the model's bias, its variance, and irreducible noise. Bias refers to the error introduced by approximating a real, often complicated, relationship with a simplified model. Variance refers to how much the model's predictions would change if it were trained on a different sample drawn from the same population.

A model with high bias and low variance, such as a simple linear regression applied to a nonlinear relationship, will consistently miss the true pattern in the same way regardless of which training sample it sees; this is called underfitting. A model with low bias and high variance, such as a deep decision tree with no depth limit, will fit the training data almost perfectly but will produce wildly different predictions depending on the specific sample it was trained on; this is called overfitting. The practical craft of machine learning, choosing model complexity, regularization strength, and amount of training data, is in large part an exercise in managing this tradeoff. Techniques such as L1 and L2 regularization, cross-validation, and dropout in neural networks are not arbitrary engineering tricks; they are direct, mathematically motivated responses to the bias-variance decomposition.




Figure 2. Conceptual illustration of the bias-variance tradeoff. As model complexity increases, bias falls while variance rises; total expected error is minimized at an intermediate complexity level rather than at either extreme.

To make this decomposition concrete rather than purely theoretical, the following small simulation estimates bias and variance directly. Two hundred to five hundred independent training samples of ten points are drawn from the relationship y equals x squared plus Gaussian noise, and a simple linear regression is fit to each sample. Because the true relationship is quadratic but the model is linear, every fitted line should miss the curvature in a similar, systematic way, while the specific slope and intercept will vary somewhat from sample to sample. Figure 3 shows the code used to run this experiment and its output, computed directly rather than assumed.

```

import numpy as np
from sklearn.linear_model import LinearRegression

# Simulate 500 resamples of n=10 points from y = x^2 + noise
rng = np.random.default_rng(0)
x_grid = np.linspace(-1, 1, 50)
preds = []

for _ in range(500):
    x = rng.uniform(-1, 1, 10)
    y = x**2 + rng.normal(0, 0.1, 10)
    model = LinearRegression().fit(x.reshape(-1,1), y)
    preds.append(model.predict(x_grid.reshape(-1,1)))

preds = np.array(preds)
bias_sq = (preds.mean(axis=0) - x_grid**2)**2
variance = preds.var(axis=0)

print("Mean bias^2:", bias_sq.mean().round(4))
print("Mean variance:", variance.mean().round(4))

# Output:
# Mean bias^2: 0.0982
# Mean variance: 0.0337

```

Figure 3. Python code and actual output for a bias-variance simulation. Five hundred independent samples of ten points are drawn from $y = x^2 + \text{noise}$, and a linear model is fit to each. The squared bias (0.0982) substantially exceeds the variance (0.0337), confirming that an underpowered linear model on a quadratic relationship is dominated by bias rather than variance, as the theory in this section predicts.

This same logic extends, with an important caveat, to deep learning. A neural network with many more parameters than training examples can easily enter a high-variance regime, prone to memorizing idiosyncrasies of its specific training set rather than learning the general pattern, as classical bias-variance intuition would predict. Figure 4 illustrates this phenomenon directly: training loss continues to fall as a model trains for more epochs, but validation loss, the error measured on data the model has not seen, eventually reaches a minimum and rises again, marking the point at which the model has stopped learning generalizable structure and started memorizing the training sample's idiosyncrasies, namely its variance.

The classical picture is not the whole story for very large models, however. Empirical work on what has been termed the double descent phenomenon (Belkin et al., 2019) has shown that, when model capacity is increased far beyond the point of perfectly fitting the training data, test error can decrease again after an initial rise, rather than continuing to climb as the simple bias-variance picture would suggest. This means the blanket statement that bigger neural networks are simply higher variance is an oversimplification: very large, heavily

overparameterized models trained with modern techniques do not always behave like the classical high-variance regime of Figure 2, and the relationship between parameter count and generalization error in that regime remains an active area of statistical research rather than a fully settled question.

Modern deep learning practice manages overfitting risk, where it does occur, through a combination of statistical tools: large amounts of training data to reduce variance directly, regularization penalties that constrain parameter magnitudes, dropout layers that inject noise during training to discourage reliance on any single feature combination, and early stopping rules that halt training once validation error, rather than training error, stops improving, as illustrated in Figure 4. Each of these is a direct, traceable application of the bias-variance framework rather than an unrelated engineering convention, even if that framework requires the double-descent caveat above to fully describe very large models.




Figure 4. Illustrative training and validation loss curves over fifty epochs. Training loss decreases monotonically, but validation loss reaches a minimum around epoch 25 and rises sharply thereafter, even as training loss keeps falling, marking the onset of overfitting and the point at which early stopping would be applied.

2.3 Sampling, Representativeness, and the Population Problem

A model can only be as good as the statistical relationship between its training sample and the population it will eventually be used on. Representative sampling, typically achieved through random or stratified sampling, is one of the most common and reliable ways to obtain

an unbiased estimator of a population quantity, although it is not the only route; estimators can also be unbiased under other designs, such as known, correctable sampling probabilities, even when the raw sample itself is not representative. What matters formally is whether the sampling mechanism is accounted for in how the estimator is constructed, not representativeness as such. In machine learning, however, training data is rarely collected through any designed sampling procedure where such corrections could be applied. Instead, datasets are frequently assembled from whatever historical records, web-scraped content, or convenience samples are available, with no accounting for how the data-generating process may have favored some parts of the population over others.

This matters statistically because a model trained on a non-representative sample will produce a biased estimator, not in the colloquial sense of being unfair, but in the precise statistical sense that its expected output differs systematically from the true population value. If a sample underrepresents a subgroup, the model's loss function, which is typically an average error over the training sample, will place very little weight on getting predictions right for that subgroup. The model is not “choosing” to be inaccurate for an underrepresented group; it is behaving exactly as the mathematics of expected loss minimization predicts.

It is worth being precise about the statistical vocabulary here, since the popular press often uses the word “bias” loosely. In statistics, an estimator is called biased if its expected value, averaged over repeated sampling, differs from the true parameter it is meant to estimate. A training procedure that minimizes average loss over a sample is, formally, computing a sample estimate of the population risk. If the sample distribution $P\text{-hat}$ differs from the true deployment population P , then minimizing risk under $P\text{-hat}$ does not generally minimize risk under P , and the size of that gap is itself a function of how different the two distributions are. This phenomenon, sometimes called distribution shift or dataset shift in the machine learning literature, is a direct descendant of classical statistical concerns about external validity in survey sampling and experimental design.

Several distinct mechanisms can produce this kind of shift. Selection bias occurs when the process used to collect the training data systematically favors certain subpopulations, for instance when a medical dataset is drawn only from patients at a single urban hospital rather than from a broader, more representative population. Measurement bias occurs when the outcome variable itself is recorded differently across groups, for instance when arrest records are used as a proxy for criminal behavior despite arrest rates being influenced by policing intensity that varies by neighborhood. Temporal bias occurs when a model trained on historical

data is deployed in a present-day context that has since shifted, a problem familiar to anyone who has watched a time series forecasting model perform well in backtesting and poorly once deployed. Each of these mechanisms can be described with the same underlying statistical language, even though they arise from very different real-world data collection processes.

3. Case Studies: When Statistical Bias Becomes Social Bias

3.1 Hiring Algorithms

Several major technology companies have experimented with automated systems to screen job applicant resumes. One widely reported case involved an internal recruiting tool built by Amazon, which was trained on a decade of historical hiring decisions (Dastin, 2018). Because the technical workforce in the training data was predominantly male, the model learned statistical associations between certain words, including the word “women’s” as in “women’s chess club,” and lower hiring scores. From a purely statistical standpoint, the model correctly identified a pattern in its training sample; the pattern itself reflected historical hiring bias rather than any genuine relationship between gender and job performance. The system was ultimately abandoned once these patterns were discovered.

This case illustrates a key statistical point: a model is an estimator of associations present in its training sample, and it has no mechanism for distinguishing between a statistically valid signal and a historically contingent artifact. Removing demographic variables from the input features does not necessarily fix the issue, because correlated proxy variables, such as participation in particular extracurricular activities, can still encode the same information, a problem statisticians refer to as omitted variable bias combined with proxy discrimination.

A useful statistical exercise is to ask what an unbiased version of this same model would have required. If the goal is to predict job performance, the correct population to sample from is, in principle, the population of all people who could plausibly perform the job well, not merely the population of people the company has previously chosen to hire. Historical hiring records are therefore a biased sample with respect to that target population by construction, since they reflect prior decisions that were themselves shaped by factors unrelated to job performance. This is sometimes called label bias in the machine learning fairness literature (Mehrabi et al., 2021): the outcome variable being predicted is already contaminated by the same process the model is meant to improve upon, which means no amount of additional training data from the same biased process can fix the underlying statistical problem.

3.2 Facial Recognition and Subgroup Error Rates

Research evaluating commercial facial recognition and gender classification systems has found substantial disparities in error rates across demographic subgroups. The Gender Shades study found that classification error rates were dramatically higher for darker-skinned women than for lighter-skinned men across several commercial systems (Buolamwini & Gebru, 2018). The statistical explanation traces back to training data composition: datasets historically used to train and benchmark these systems have been disproportionately composed of lighter-skinned individuals, particularly men.

This is a direct illustration of the sampling representativeness problem described in Section 2.3. An algorithm minimizing average error over an unrepresentative sample will, by construction, achieve low average error while exhibiting high error specifically on the underrepresented subgroup, since that subgroup contributes little to the overall average being minimized. The aggregate accuracy metric, which is often the only number reported, can look excellent while masking severe and statistically predictable disparities in subgroup performance.




Figure 5. Illustrative pattern of subgroup error rate disparity reported in audits of commercial gender classification systems, with substantially higher error rates for darker-skinned women than for lighter-skinned men. Values are rounded and illustrative of the general pattern reported in the literature (Buolamwini & Gebru, 2018) rather than figures from a single specific system.

3.3 Criminal Risk Assessment Tools

Statistical risk assessment tools used in some criminal justice settings to predict the likelihood that a defendant will reoffend have been the subject of extensive statistical scrutiny. A ProPublica analysis of the COMPAS risk assessment tool found that it produced higher false

positive rates, incorrectly flagging individuals as high risk, for Black defendants compared with white defendants (Angwin et al., 2016), while the tool's developers countered that the tool was well-calibrated, meaning that within each risk score category, the actual reoffense rate was similar across racial groups.

This controversy is, at its heart, a statistical one. Chouldechova (2017) and Kleinberg, Mullainathan, and Raghavan (2017) have since shown mathematically that, except in trivial cases, it is impossible for a risk score to simultaneously achieve equal calibration across groups and equal false positive and false negative rates across groups when the base rates of the outcome differ between those groups. This is not a failure of any particular algorithm's design; it is a theorem about the structure of probability and classification under unequal base rates.




Figure 6. A worked numerical illustration of the calibration / equal-error-rate tradeoff. The left panel shows a risk score that is equally calibrated for both groups: within each risk bin, the observed outcome rate is identical. The right panel shows that, because Group A's overall base rate (0.50) is higher than Group B's (0.20), this same calibrated score necessarily produces a higher false positive rate for Group A. No recalibration can equalize both properties simultaneously while base rates differ.

The case demonstrates that some fairness tradeoffs are not engineering problems waiting for a clever fix but mathematical impossibility results that require value judgments about which statistical property to prioritize.

4. Statistical Remedies and Their Limits

4.1 Reweighting and Resampling

One of the most direct statistical responses to an unrepresentative training sample is to reweight or resample the data so that underrepresented groups contribute proportionally more to

the loss function being minimized. Inverse propensity weighting, a technique borrowed from survey statistics and causal inference, assigns larger weights to observations from groups that were undersampled relative to their true population share, effectively correcting the loss function to approximate what it would have been under representative sampling. Oversampling minority classes, or generating synthetic examples for them through methods such as the synthetic minority oversampling technique (Chawla et al., 2002), serves a similar statistical purpose for classification problems with rare outcome categories.

These techniques are effective but not magical. Reweighting can reduce the variance penalty associated with a small subgroup sample only up to a point; if the actual number of distinct observations for that subgroup is small, no amount of reweighting manufactures new information, and the resulting estimator, while less biased on average, may still have unacceptably high variance. This is a direct consequence of standard statistical results on the variance of a weighted estimator, which depends on the effective sample size rather than the nominal sample size after reweighting.

4.2 Fairness Constraints and Their Statistical Tradeoffs

A second family of remedies imposes explicit fairness constraints during model training, for example requiring that the model's false positive rate be approximately equal across demographic groups, or that its positive prediction rate be approximately equal across groups. These constraints can be incorporated directly into the optimization objective alongside the usual loss function, producing a constrained estimation problem rather than an unconstrained one.

As discussed in Section 3.3, however, statistical theory shows that several intuitively desirable fairness criteria cannot generally be satisfied simultaneously when outcome base rates differ across groups. Choosing which fairness constraint to impose is therefore not a purely technical decision; it encodes a value judgment about which type of error, a false positive or a false negative, is more costly to avoid for a given application, and for whom. A statistically literate practitioner can make this tradeoff explicit and quantifiable, for instance by reporting how much overall predictive accuracy is sacrificed to achieve a given reduction in a specific disparity metric, rather than presenting a single fairness intervention as though it resolves the issue completely.

4.3 The Limits of Post Hoc Correction

A final point follows directly from basic estimation theory: no amount of post hoc statistical correction can substitute for the information that a poorly designed data collection process failed to capture. If a particular subpopulation was never measured under a relevant condition, for instance a medical sensor calibrated and validated almost entirely on one skin tone range, then no reweighting, resampling, or constraint can recover the missing information; it can only redistribute the existing uncertainty differently across groups. This is why many statisticians and machine learning researchers now argue that the most durable solution to many fairness problems is improving data collection design itself, including deliberate oversampling of underrepresented groups during initial data gathering, rather than relying exclusively on algorithmic correction applied after the fact.

4.4 Statistical Auditing in Practice

Beyond remedies applied during training, a complementary statistical practice is auditing a model after it has been deployed. A statistical audit typically involves stratifying model performance by relevant subgroups and reporting separate accuracy, calibration, and error rate metrics for each stratum, rather than a single aggregate number. This is directly analogous to checking whether the residuals of a regression model are randomly scattered or instead show a systematic pattern when plotted against a covariate; in both cases, an aggregate summary statistic can hide a structured failure that only becomes visible once the data are disaggregated.

A complete statistical audit also reports confidence intervals around subgroup performance metrics rather than point estimates alone. This matters because smaller subgroups naturally produce noisier estimates: an observed difference in error rate between a large group and a small group may reflect genuine disparity, or it may simply reflect the higher sampling variability of the smaller group's estimate. Distinguishing between these two explanations requires the same statistical hypothesis testing machinery taught in an introductory statistics course, applied to model outputs rather than to a survey or experiment. This observation reinforces the broader theme of this paper: rigorous evaluation of AI systems is, in large part, an exercise in applied statistics.

5. Discussion: Why These Are Statistical Problems, Not Just Engineering Problems

A recurring theme across the case studies in Section 3 is that the failures described are not random bugs but predictable outcomes of well-understood statistical principles operating on imperfect data. This distinction matters for how the problems should be addressed. An

engineering bug, such as a software crash, can typically be fixed by correcting the faulty code. A statistical bias arising from unrepresentative sampling or from a mathematical impossibility result cannot be fixed by writing better code alone; it requires either collecting better data, accepting an explicit tradeoff between competing statistical fairness criteria, or constraining the model with additional statistical assumptions informed by domain knowledge.

This reframing has practical implications for math and computer science students entering the field. A student who understands the bias-variance tradeoff, the mechanics of maximum likelihood estimation, and the conditions under which sample statistics generalize to populations is better equipped to anticipate where a deployed model will fail than a student who only understands how to call a machine learning library's ".fit()" function. The statistical perspective also clarifies that model auditing, the practice of testing a deployed model's performance across subgroups and conditions, is not an optional add-on but a natural extension of standard statistical validation practice, akin to checking the residuals of a regression model.

It is also worth noting the limits of statistical correction. Reweighting a training sample to be more representative, a common bias mitigation technique, can reduce but not always eliminate disparities, particularly when the available data for an underrepresented group is small enough that even an unbiased estimator computed from it will have high variance. In such cases, more data collection, rather than algorithmic adjustment alone, is often the more durable statistical solution.

Finally, the statistical framing offered in this paper suggests a practical checklist for evaluating any deployed predictive system, whether built by a student for a class project or by a large technology company for production use. The relevant questions include the following: What population was the training sample drawn from, and how does it compare to the population the model will actually be applied to? Is overall accuracy being reported, or is performance broken out by relevant subgroups where disparities could otherwise hide inside an average? Are the model's probability outputs calibrated, or merely useful for ranking, as discussed in Section 2.1 and Figure 1? And when a fairness intervention is applied, has its statistical cost, in terms of accuracy or other competing fairness criteria, been made explicit rather than assumed away? None of these questions require advanced machine learning expertise to ask; they require statistical literacy, which is precisely the point this paper has tried to establish.

6. Conclusion

This paper has argued that the behavior of machine learning systems, including their failures, is best understood through the lens of classical statistical theory rather than treated purely as a software engineering concern. The bias-variance tradeoff, illustrated both conceptually in Figure 2 and computationally in Figure 3, explains the fundamental tension between underfitting and overfitting that governs model design choices. Sampling theory explains why models trained on unrepresentative data produce systematically biased predictions for underrepresented groups. Case studies in hiring, facial recognition, and criminal risk assessment demonstrate that these statistical phenomena have measurable, real-world consequences, and that some apparent fairness failures, as Figure 6 makes concrete, are in fact mathematically provable tradeoffs rather than fixable bugs.

For students pursuing mathematics and computer science, the implication is clear: a rigorous grounding in probability and statistics is not a prerequisite to be completed before “real” machine learning begins, but an ongoing analytical tool needed to build, evaluate, and responsibly deploy predictive systems. As AI systems continue to expand into high-stakes domains such as healthcare, lending, and criminal justice, the capacity to reason statistically about bias, uncertainty, and the limits of estimation will only grow in importance.

More broadly, this paper suggests that the conversation around artificial intelligence and fairness would benefit from greater fluency in the statistical vocabulary already developed over the past century: bias and variance, calibration and discrimination, selection effects and base rates. These are not new concepts invented to critique AI; they are the same tools statisticians have used for decades to evaluate surveys, clinical trials, and economic forecasts. Applying that existing vocabulary carefully to machine learning systems does not resolve every disagreement about what fairness should mean in a given context, but it does ensure that the disagreement is about values and priorities rather than about confusion over what the underlying mathematics actually implies. For a student of mathematics and computer science, that clarity is itself a valuable and transferable skill, applicable far beyond any single model or dataset.

References

- Angwin, J., Larson, J., Mattu, S., & Kirchner, L. (2016). Machine bias: There's software used across the country to predict future criminals. *ProPublica*.
- Belkin, M., Hsu, D., Ma, S., & Mandal, S. (2019). Reconciling modern machine-learning practice and the classical bias-variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32), 15849-15854.
- Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. *Proceedings of Machine Learning Research*, 81, 77-91.
- Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16, 321-357.
- Chouldechova, A. (2017). Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2), 153-163.
- Dastin, J. (2018). Amazon scraps secret AI recruiting tool that showed bias against women. *Reuters*.
- Dawid, A. P. (1982). The well-calibrated Bayesian. *Journal of the American Statistical Association*, 77(379), 605-610.
- Friedman, J., Hastie, T., & Tibshirani, R. (2009). *The elements of statistical learning: Data mining, inference, and prediction* (2nd ed.). Springer.
- Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017). On calibration of modern neural networks. *Proceedings of the 34th International Conference on Machine Learning*, 70, 1321-1330.
- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). *An introduction to statistical learning: With applications in R* (2nd ed.). Springer.
- Kleinberg, J., Mullainathan, S., & Raghavan, M. (2017). Inherent trade-offs in the fair determination of risk scores. *Proceedings of Innovations in Theoretical Computer Science*.
- Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., & Galstyan, A. (2021). A survey on bias and fairness in machine learning. *ACM Computing Surveys*, 54(6), 1-35.
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.